

THE PLAN

Your first useful AI workflow.

A 30-day working kit to remove a bottleneck, support the people doing the work, and make the next step worth taking.

A month to learn. Not a promise to launch.

Use this with a business owner, someone who does the work, and IT or security support. Choose a bounded task. A regulated, high-impact or complex workflow may need a longer review before testing.

The four conversations

Days 1-7: Find the work. Map one recurring problem and what better would mean.

Days 8-14: Set the boundaries. Agree approved information, actions, human review and a fallback.

Days 15-21: Practice and observe. Test with a small group and representative, permitted examples.

Days 22-30: Review the evidence. Compare quality, effort and full cost. Continue, change or stop.

Business owner / team / review date

--

Keep completed worksheets inside your approved environment. Use a PDF reader that supports forms, then save a local copy. No entries are sent to AIR by this PDF.

DAYS 1-7

Choose work worth improving.

Start with a recurring task that is understandable, measurable and reversible. A good first test leaves a person able to check the result.

Example: prepare a draft explanation of a spending variance from an approved, de-identified sample. A finance owner verifies the source and interpretation. The tool does not approve spend or alter accounting records.

What is the task, who does it, and where does it stall?

Today: time, frequency, errors/rework, and the effect on people

What improvement would count? What must not get worse?

Could a process change or a capability we already own solve this?

DAYS 8-14

Give the work boundaries.

This is an operating worksheet, not an approved policy. Have the relevant business, security, privacy and legal owners review the proposed use before a pilot.

Approved tool/environment, owner, and information it may use

Information and actions excluded from this pilot

What needs human review? Who may approve consequential actions?

When do we pause or hand over? Who responds, and what is the fallback?

If an agent can act: list each system it may reach, permissions, spending/action limits, activity records and who can revoke access. Test those limits before connecting it to live work.

DAYS 15-21

Make room for people to learn.

Ask the team what feels difficult, what should stay in their hands, and what a useful improvement would free them to do. Give them time to practice and a person to ask.

A reusable instruction for an approved LLM

“Help me draft [output] for [audience] using only [approved source]. The purpose is [business need]. Follow [format and limits]. Do not invent missing facts. Identify assumptions, missing information and source passages for material claims. Ask questions if the task is unclear. Produce a draft for human review; do not take action.”

A clear prompt helps define the task; it does not make an answer correct. Verify source references and claims. Do not put sensitive information into an unapproved service. Product features and data settings must be checked for the account you actually use.

Two representative examples and one difficult exception to test

Who reviews the output? What errors or extra work did they find?

Team feedback / time to learn / intended use of time saved

BEFORE EXPANDING

Look beyond the starting price.

Compare the same workload across your current approach, existing tools, and any proposed purchase. Record assumptions rather than inventing a return on investment.

One-time costs: setup, integration, migration, training and testing

Ongoing costs: licenses/usage, data transfer, review, support and maintenance

Expected and higher-demand case: volume, response time and capacity

Commitment and exit: minimums, renewal, data export, switching effort

Time saved is not automatically cash saved. Explain whether capacity is redeployed, service improves, overtime falls, or a real expense changes. Avoid counting the same benefit twice. Ask an advisor to help resolve material supplier or infrastructure unknowns.

DAYS 22-30

Bring back a useful answer.

The result can be “continue,” “change the approach,” or “not yet.” The purpose of the pilot is evidence, not defending the purchase.

What improved? What did quality, effort, cost and team feedback show?

Continue / change / stop, with the reason and unresolved risks

Next action / accountable owner / evidence needed / review date

Bring the worksheet into the conversation.

AIR helps connect business priorities with technology options, specialists and suppliers. You remain in charge. If a data, network, security or supplier question is blocking the next step, bring that question with you.

Visit www.airesulting.com to learn about working with AI Resulting.

Research context: This original AIR worksheet responds to themes of workflow value, governance and workforce support in McKinsey, *State of AI* (August 25, 2026), RSM, *Middle Market AI Survey 2026*, and Deloitte, *State of AI* (January 21, 2026). It is not their assessment or an endorsed methodology.