

A working kit

Your first useful *AI workflow.*

Take one recurring task from bottleneck to evidence.

Work through it with the business owner, a person who does the task, and IT or security support. Pick something bounded. Regulated or high-impact work may need a longer review before anyone tests it.

Four conversations. Set a review date before the first one.

Find the work.	Map one recurring problem and what better would look like.
Set the boundaries.	Agree what the tool may use and do, who reviews it, and what it depends on underneath.
Practice and observe.	Test with a small group on permitted examples.
Review the evidence.	Compare quality, effort and full cost. Then decide: continue, change or stop.

Business owner, team and review date

A planning aid for your own team. Security, privacy and legal review still sit with your owners. Type into it with any PDF reader that supports forms and save a copy. Nothing you type is sent anywhere.

Prefer the browser? The same worksheet is at airesulting.com/ai-workflow-worksheet

Find the work

Choose work worth improving.

Pick a recurring task you can measure and undo if it goes wrong. A person should be able to check every result.

Example. Draft an explanation of a spending variance from an approved, de-identified sample. A finance owner verifies the source and the interpretation. The tool doesn't approve spend or change accounting records.

What is the task, who does it, and where does it stall?

How it runs today: time, frequency, errors or rework, and the effect on people

What improvement would count? What must not get worse?

Could a process change or a tool you already own solve this?

Set the boundaries

Give the work boundaries.

This worksheet is an operating plan, not an approved policy. Have your business, security, privacy and legal owners review the proposed use before a pilot.

Approved tool or environment, its owner, and the information it may use

Information and actions excluded from this pilot

What needs human review? Who may approve consequential actions?

When do you pause or hand over? Who responds, and what's the fallback?

If an agent can take actions, list each system it can reach, its permissions and spending limits, where its activity is recorded and who can revoke access. Test those limits before it touches live work.

Set the boundaries

What does it depend on underneath?

The tool is the visible part. Whether it works depends on what sits under it. Answer these before the pilot, while a gap is still cheap to fix.

Systems and data

Which systems and data does the task need? Are they connected, current and trusted?

Access and permissions

Who can reach that information today? Do the permissions still hold once the tool can reach it too?

Network and environment

Can your connectivity, devices and hosting carry the added work, including at busy times?

Security and governance

What security or governance review does this need, and who signs off?

If the honest answer is “we don’t know,” write that down. It’s the most useful answer on this page.

Practice and observe

Make room for people to learn.

Ask the team what feels hard, what should stay in their hands, and what a real improvement would free them to do. Give them time to practice and someone to ask.

A reusable instruction for an approved LLM

“Help me draft [output] for [audience] using only [approved source]. The purpose is [business need]. Follow [format and limits]. Do not invent missing facts. Identify assumptions, missing information and source passages for material claims. Ask questions if the task is unclear. Produce a draft for human review; do not take action.”

Check the sources and claims before anyone relies on a draft. Keep sensitive information out of unapproved services.

Two representative examples and one difficult exception to test

Who reviews the output? What errors or extra work did they find?

Team feedback, time to learn, and where the time saved will go

Review the evidence

Count the full cost.

Compare the same workload across your current approach, tools you already own, and any proposed purchase. Record assumptions instead of inventing a return.

One-time costs: setup, integration, migration, training and testing

Ongoing costs: licenses or usage, data transfer, review, support and maintenance

Expected and higher-demand cases: volume, response time and capacity

Commitment and exit: minimums, renewal, data export and switching effort

Time saved is not automatically cash saved.

Say whether capacity is redeployed, service improves, overtime falls or a real expense changes. Count each benefit once.

Review the evidence

Bring back a useful answer.

Continue, change the approach, or not yet. Any of the three is a good result if the evidence supports it.

What improved? What did quality, effort, cost and team feedback show?

Continue, change or stop, with the reason and the risks still open

Next action, accountable owner, evidence needed and review date

Take it to a Resultant.

Your answers are the agenda for a first conversation. We'll look at what the next step depends on and help you line up the right people and options to move it forward.

[Talk to a Resultant at airesulting.com/talk](https://airesulting.com/talk) ↗

Research context: this original AIR worksheet responds to themes of workflow value, governance and workforce support in [McKinsey, The State of AI \(August 25, 2026\)](#); [RSM, Middle Market AI Survey 2026](#); and [Deloitte, State of AI \(January 21, 2026\)](#). It is not their assessment or an endorsed methodology.